

Filesystems


COSC-361
Stephen Marz



MIN H. KAO DEPARTMENT OF
ELECTRICAL ENGINEERING &
COMPUTER SCIENCE

1

Data Route



2

System Calls

- open
- read
- write
- lseek
- close

These operate on a "generic" file.

3

Virtual File System

- This system is in the kernel
 - It virtualizes the file system from the user application.
- It allows us to have a unified file tree that can support a multitude of file systems.
 - Example: /home on Hydra is an NFS drive
- We use the same system calls regardless of underlying file system type, thanks to VFS.

4 3/24/2019

COSC 361



4

Block Device Drivers

```

struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
    ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
    int (*iterate) (struct file *, struct dir_context *);
    int (*iterate_shared) (struct file *, struct dir_context *);
    __poll_t (*poll) (struct file *, struct poll_table_struct *);
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    unsigned long mmap_supported_flags;
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, loff_t, loff_t, int datasync);
}

```

Function pointers are used for every file operation. This way, all I do is provide a function pointer that handles the given operation (open, close, etc).

5 3/24/2019

COSC 361



5

Disk Protocols

- (S)ATA [S for Serial vs Parallel]
 - Advanced Technology Attachment
 - AT is for the IBM AT
- ATAPI
 - ATA Packet Interface (typically for CDROMs)
- SCSI
 - Small computer system interface
- SAS
 - Serial attached SCSI
- NVME
 - Non-volatile memory express
 - For PCI-express SSDs

6 3/24/2019

COSC 361



6

Input / Output

- PCI
 - Peripheral Component Interconnect
- IDE
 - Integrated Drive Electronics
- MMIO
 - Memory-mapped I/O


7 3/24/2019

COSC 361



7

Disks



8 3/24/2019

COSC 361



8

File Systems

- Unix-style
 - Superblock
 - Inodes
 - Blocks

9 3/24/2019

COSC 361



9

Superblock

```

struct super_operations {
    struct inode *alloc_inode(struct super_block *sb);
    void (*destroy_inode)(struct inode *);

    void (*dirty_inode)(struct inode *, int flags);
    int (*write_inode)(struct inode *, struct writeback_control *wbc);
    int (*drop_inode)(struct inode *);
    void (*wait_inode)(struct inode *);
    void (*wait_super)(struct super_block *);
    int (*sync_fs)(struct super_block *sb, int wait);
    int (*freeze_super)(struct super_block *);
    int (*unfreeze_fs)(struct super_block *);
    int (*show_super)(struct super_block *);
    int (*setfs)(struct dentry *, struct kstatfs *);
    int (*remount_fs)(struct super_block *, int *, char *);
    void (*umount_begin)(struct super_block *);

    int (*show_options)(struct seq_file *, struct dentry *);
    int (*show_devname)(struct seq_file *, struct dentry *);
    int (*show_path)(struct seq_file *, struct dentry *);
    int (*show_stats)(struct seq_file *, struct dentry *);

#ifdef CONFIG_QUOTA
    ssize_t (*quota_read)(struct super_block *, int, char *, size_t, loff_t);
    ssize_t (*quota_write)(struct super_block *, int, const char *, size_t, loff_t);
    struct kstatfs (*get_kstatfs)(struct inode *);
#endif
    int (*bdev_try_to_free_page)(struct super_block *, struct page *, gfp_t);
    long (*nr_cached_objects)(struct super_block *,
        struct shrink_control *);
    long (*free_cached_objects)(struct super_block *,
        struct shrink_control *);
};

```

10 3/24/2019

COSC 361



10

Inode

```

struct inode_operations {
    struct dentry * (*lookup)(struct inode *, struct dentry *, unsigned int);
    const char * (*get_link)(struct dentry *, struct inode *, struct delayed_call *);
    int (*permission)(struct inode *, int);
    struct posix_acl * (*get_acl)(struct inode *, int);

    int (*readlink)(struct dentry *, char __user *, int);

    int (*create)(struct inode *, struct dentry *, umode_t, bool);
    int (*link)(struct inode *, struct inode *, struct dentry *);
    int (*unlink)(struct inode *, struct dentry *);
    int (*symlink)(struct inode *, struct dentry *, const char *);
    int (*mkdir)(struct inode *, struct dentry *, umode_t);
    int (*rmdir)(struct inode *, struct dentry *);
    int (*mknod)(struct inode *, struct dentry *, umode_t, dev_t);
    int (*rename)(struct inode *, struct dentry *,
        struct inode *, struct dentry *, unsigned int);
    int (*setattr)(struct dentry *, struct iattr *);
    int (*getattr)(const struct path *, struct kstat *, u32, unsigned int);
    ssize_t (*listxattr)(struct dentry *, char *, size_t);
    int (*fiemap)(struct inode *, struct fiemap_extent_info *, u64 start,
        u64 len);
    int (*update_time)(struct inode *, struct timespec64 *, int);
    int (*atomic_open)(struct inode *, struct dentry *,
        struct file *, unsigned open_flag,
        umode_t create_mode, int *opened);
    int (*tmpfile)(struct inode *, struct dentry *, umode_t);
    int (*set_acl)(struct inode *, struct posix_acl *, int);
    __cacheline_aligned
};

```

11 3/24/2019

COSC 361



11

Blocks

- Blocks generally cannot be contiguous if files can be expanded at any time.
- Blocks are stored in "chunks" on the hard drive.
 - Typical chunk sizes:
 - 4096 bytes
 - 8192 bytes
- Every file must take a multiple of these chunk size. 1 byte file still takes 4096 bytes for 1 block.

12 3/24/2019

COSC 361



12

Locating Blocks

- The inode contains pointers that point to the blocks.
 - This process is known as *indirection*.
 - This is necessary since inode is a fixed size and you can have a large number of blocks.

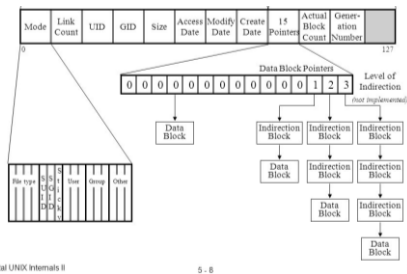
13

COSC 361



13

Block Indirection (UFS)



14

COSC 361



14

File Modes

- Describe permissions and other metadata for a file.
 - S_IFDIR - Directory bit
 - S_IFCHR - Char device bit
 - S_IFBLK - Block device bit
 - S_IFIFO - FIFO device bit

15

COSC 361



15

Octal Modes

- Modes are specified in octal (base 8)
- 0777 = 0b000_111_111_111
- Octal specifies groups of 3 bits
 - Specified in four groups of 3
 - Group 3 [leftmost] = metadata (DIR/CHR/BLK/REG/FIFO bits)
 - Group 2 = Owner bits Read, Write, and Execute
 - Group 1 = Group bits Read, Write, and Execute
 - Group 0 = Other bits Read, Write, and Execute

16 3/24/2019

COSC 361



16

Example

- I want a regular file with:
 - Owner = read and write
 - Group = read only
 - Other = execute only

0b000_110_100_001 = 0641

000_RWX_RWX_RWX

OWN_GRP_OTH

17 3/24/2019

COSC 361



17

Combined Lists

- RWX permissions
 - chmod = change mode
- Typically shown as a string:
 - -rwxrwxrwx
 - -r-Xr---WX
 - In this case
 - the owner can read and execute
 - the group can read only
 - any others can write and execute

18 3/24/2019

COSC 361



18

Access Control Lists

- An exhaustive list of who can do what to a file.
 - Much more access control
 - Much less efficient
 - Cannot be implemented as a 16-bit integer

19

COSC 361



19

Filesystems

COSC-361
Stephen Marz



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

20